

and there many other ways of encoding text that could be explored. Finally, the models used here are relatively simple, and there are other architectures that could have been used such as RNNs, or autoencoders, that are more complex and may perform better.